

linear_models_ex_02

June 24, 2024

1 Exercise M4.02

In the previous notebook, we showed that we can add new features based on the original feature `x` to make the model more expressive, for instance `x ** 2` or `x ** 3`. In that case we only used a single feature in `data`.

The aim of this notebook is to train a linear regression algorithm on a dataset with more than a single feature. In such a “multi-dimensional” feature space we can derive new features of the form `x1 * x2`, `x2 * x3`, etc. Products of features are usually called “non-linear” or “multiplicative” interactions between features.

Feature engineering can be an important step of a model pipeline as long as the new features are expected to be predictive. For instance, think of a classification model to decide if a patient has risk of developing a heart disease. This would depend on the patient’s Body Mass Index which is defined as `weight / height ** 2`.

We load the dataset penguins dataset. We first use a set of 3 numerical features to predict the target, i.e. the body mass of the penguin.

Note

If you want a deeper overview regarding this dataset, you can refer to the Appendix - Datasets description section at the end of this MOOC.

```
[14]: import pandas as pd

penguins = pd.read_csv("../datasets/penguins.csv")

columns = ["Flipper Length (mm)", "Culmen Length (mm)", "Culmen Depth (mm)"]
target_name = "Body Mass (g)"

# Remove lines with missing values for the columns of interest
penguins_non_missing = penguins[columns + [target_name]].dropna()

data = penguins_non_missing[columns]
target = penguins_non_missing[target_name]
data.head()
```

```
[14]:   Flipper Length (mm)  Culmen Length (mm)  Culmen Depth (mm)
0                181.0                39.1                18.7
```

1	186.0	39.5	17.4
2	195.0	40.3	18.0
4	193.0	36.7	19.3
5	190.0	39.3	20.6

Now it is your turn to train a linear regression model on this dataset. First, create a linear regression model.

```
[15]: from sklearn.linear_model import LinearRegression

linear_regression = LinearRegression()
```

Execute a cross-validation with 10 folds and use the mean absolute error (MAE) as metric.

```
[16]: from sklearn.model_selection import cross_validate

results = cross_validate(linear_regression, data, target, scoring =_
    ↪ "neg_mean_absolute_error", return_train_score=True, n_jobs=2, verbose=1,_
    ↪ cv=10)
```

```
[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done 10 out of 10 | elapsed: 1.2s finished
```

Compute the mean and std of the MAE in grams (g). Remember you have to revert the sign introduced when metrics start with `neg_`, such as in `"neg_mean_absolute_error"`.

```
[17]: print(f"{-results['test_score'].mean()} +- {results['test_score'].std()}")
```

```
337.07133738443895 +- 84.86840942516221
```

Now create a pipeline using `make_pipeline` consisting of a `PolynomialFeatures` and a linear regression. Set `degree=2` and `interaction_only=True` to the feature engineering step. Remember not to include a “bias” feature (that is a constant-valued feature) to avoid introducing a redundancy with the intercept of the subsequent linear regression model.

You may want to use the `.set_output(transform="pandas")` method of the pipeline to answer the next question.

```
[18]: from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import PolynomialFeatures

poly_features = PolynomialFeatures(degree=2, interaction_only=True)
model = make_pipeline(poly_features, linear_regression).
    ↪ set_output(transform="pandas")
model
```

```
[18]: Pipeline(steps=[('polynomialfeatures',
    PolynomialFeatures(interaction_only=True)),
    ('linearregression', LinearRegression())])
```

Transform the first 5 rows of the dataset and look at the column names. How many features are generated at the output of the PolynomialFeatures step in the previous pipeline?

```
[20]: model.fit(data, target)
      model[0].transform(data[:5])
```

```
[20]:      1  Flipper Length (mm)  Culmen Length (mm)  Culmen Depth (mm)  \
0  1.0                181.0                39.1                18.7
1  1.0                186.0                39.5                17.4
2  1.0                195.0                40.3                18.0
4  1.0                193.0                36.7                19.3
5  1.0                190.0                39.3                20.6

      Flipper Length (mm) Culmen Length (mm)  \
0                        7077.1
1                        7347.0
2                        7858.5
4                        7083.1
5                        7467.0

      Flipper Length (mm) Culmen Depth (mm)  Culmen Length (mm) Culmen Depth (mm)
0                        3384.7                        731.17
1                        3236.4                        687.30
2                        3510.0                        725.40
4                        3724.9                        708.31
5                        3914.0                        809.58
```

Check that the values for the new interaction features are correct for a few of them.

```
[24]: data.iloc[0][0]*data.iloc[0][1]
```

```
[24]: 7077.1
```

Use the same cross-validation strategy as done previously to estimate the mean and std of the MAE in grams (g) for such a pipeline. Compare with the results without feature engineering.

```
[25]: results_with_p = cross_validate(model, data, target, scoring =_
    ↪ "neg_mean_absolute_error", return_train_score=True, n_jobs=2, verbose=1,_
    ↪ cv=10)
print(f"{-results['test_score'].mean()} +- {results['test_score'].std()}")
print(f"{-results_with_p['test_score'].mean()} +- {results_with_p['test_score'].
    ↪ std()}")
```

```
[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.
```

```
337.07133738443895 +- 84.86840942516221
301.78955228431437 +- 44.34000781934426
```

```
[Parallel(n_jobs=2)]: Done 10 out of 10 | elapsed: 1.2s finished
```

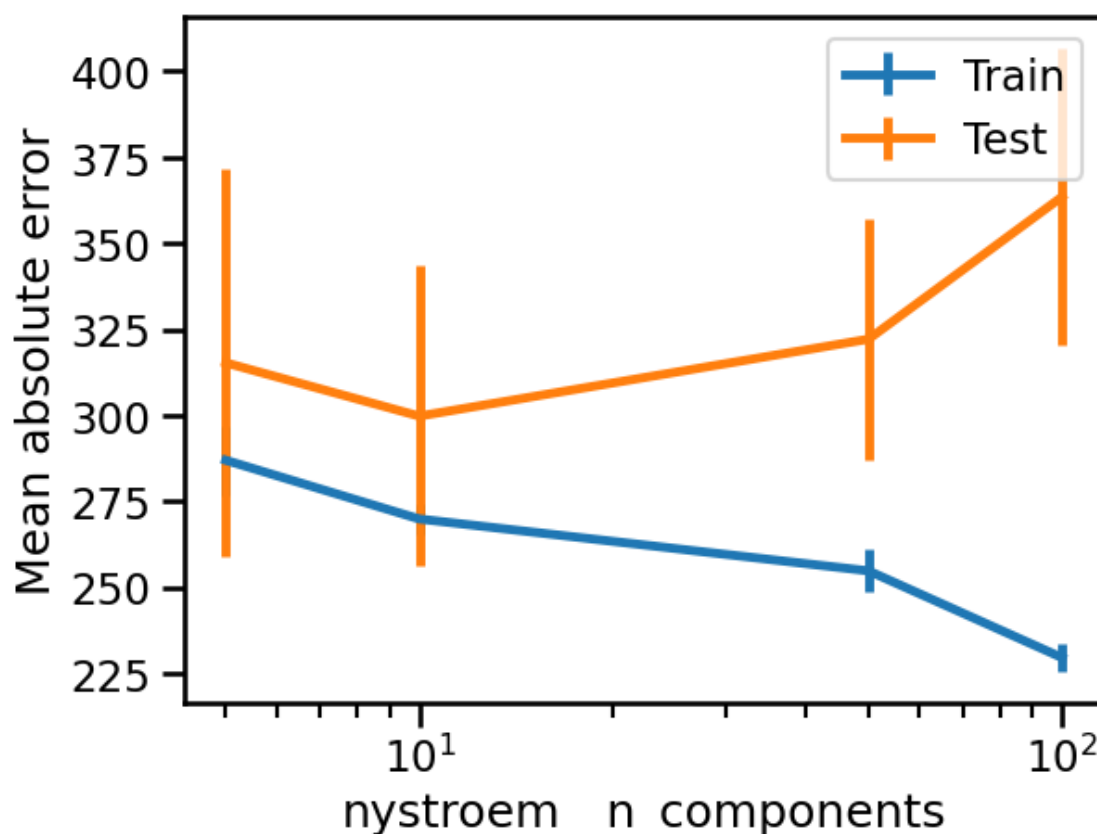
Now let's try to build an alternative pipeline with an adjustable number of intermediate features while keeping a similar predictive power. To do so, try using the `Nystroem` transformer instead of `PolynomialFeatures`. Set the `kernel` parameter to "poly" and `degree` to 2. Adjust the number of components to be as small as possible while keeping a good cross-validation performance.

Hint: Use a `ValidationCurveDisplay` with `param_range = np.array([5, 10, 50, 100])` to find the optimal `n_components`.

```
[32]: from sklearn.kernel_approximation import Nystroem
      from sklearn.model_selection import ValidationCurveDisplay
      import numpy as np

      ny = Nystroem(kernel="poly", degree=2)
      model_with_ny = make_pipeline(ny, linear_regression).
      ↪set_output(transform="pandas")
      ValidationCurveDisplay.from_estimator(model_with_ny, data, target,
      ↪param_name="nystroem_n_components", cv=10, param_range=np.array([5, 10, 50,
      ↪100]), scoring="neg_mean_absolute_error",
      negate_score=True,
      std_display_style = "errorbar",
      n_jobs =2)
```

```
[32]: <sklearn.model_selection._plot.ValidationCurveDisplay at 0x7f596c1ee910>
```



How do the mean and std of the MAE for the Nystroem pipeline with optimal `n_components` compare to the other previous models?

```
[34]: model_with_ny.set_params(nystroem__n_components=10)
cv_results = cross_validate(model, data, target, cv=10, scoring =_
    ↪ "neg_mean_absolute_error", n_jobs=2)

print(f"{-cv_results['test_score'].mean()} +- {cv_results['test_score'].std()}")
```

301.78955228431437 +- 44.34000781934426

```
[ ]:
```